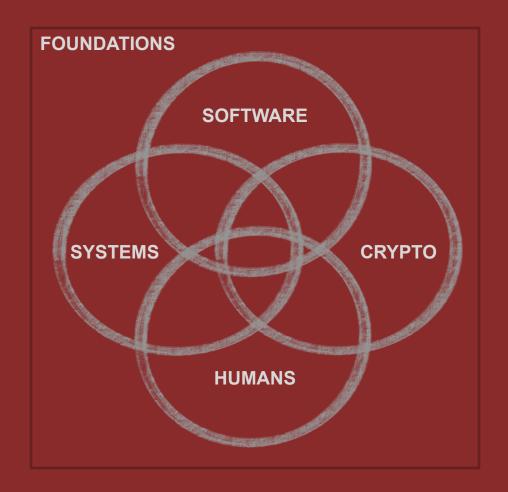
Διάλεξη #2 - x86 Basics and Buffer Overflows

Εθνικό και Καποδιστριακό Πανεπιστήμιο Αθηνών

Εισαγωγή στην Ασφάλεια

Θανάσης Αυγερινός



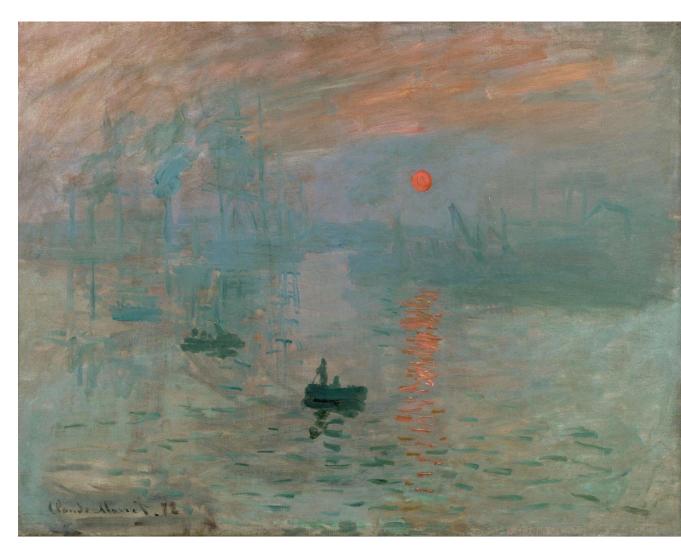
Huge thank you to <u>David Brumley</u> from Carnegie Mellon University for the guidance and content input while developing this class

Ανακοινώσεις / Διευκρινίσεις

- Μόλις ανέβηκε η εργασία #0!
 - ο Προθεσμία: 12 Μαρτίου 23:59

Την Προηγούμενη Φορά

- Security Fundamentals
 - Adversaries
 - Threat Models
 - Security Properties
 - Trusted Computing Base (TCB)
 - Security Principles



Σήμερα

- x86 Assembly Fundamentals
 - Call Return Semantics
- Basics of buffer overflow attacks



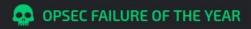
Security in the News

DISA, which provides services like drug and alcohol testing and background checks to more than 55,000 enterprises and a third of Fortune 500 companies, confirmed the data breach in a <u>filing</u> with Maine's attorney general on Monday.

DISA said it discovered it had been the victim of a "cyber incident" that affected a "limited portion" of its network on April 22, 2024. An internal investigation determined that a hacker had infiltrated the company's network on February 9, 2024, where they went unnoticed for over two months.

In a letter sent to those affected by the data breach, which includes individuals who underwent employee screening tests, DISA said the attacker "procured some information" from its systems.

In a separate <u>filing</u> with the Massachusetts attorney general, DISA confirmed the stolen information included individuals' Social Security numbers; financial account information, including credit card numbers; and government-issued identification documents. This filing confirmed that more than 360,000 Massachusetts residents were affected by the breach.



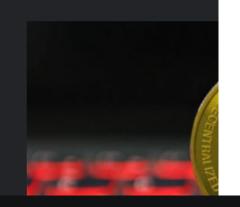
How North Korea pulled off a \$1.5 billion crypto heist—the biggest in history

Attack on Bybit didn't hack infrastructure or exploit sm code. So how did it work?

DAN GOODIN - FEB 24, 2025 6:41 PM | 134



The cryptocurrency industry and those responsible for following Friday's heist, likely by North Korea, that drai exchange Bybit, making the theft by far the biggest eve



- Operate seamlessly across Windows, MacOS, and various wallet interfaces
- Show minimal signs of compromise while maintaining persistence
- Function as backdoors to execute arbitrary commands
- Download and execute additional malicious payloads
- Manipulate what users see in their interfaces

These hackers have also been long known for their relentless social engineering prowess. They often spend weeks or months building online personas that ultimately win the trust of targets. That persistence likely allowed the thieves who hit Bybit to somehow tamper with the UIs of each company employee whose digital imprimatur was required to move the funds out of cold storage—and ultimately into wallets the hackers controlled—all at breakneck speed.

As both Check Point and Trail of Bits point out, the lessons learned here bring cryptocurrency security back to some of the most basic elements, such as segmenting internal networks, adopting defense-in-depth practices that include multiple, overlapping controls for detecting and preventing sophisticated attacks, and preparation for scenarios precisely like this one.



Reminders:
Memory, Stack,
and Endianness
(speedrunning
progintro)

Η Μνήμη Οργανώνεται σε Bytes (Υπενθύμιση)

Το μέγεθος της μνήμης μετράται σε Bytes:

- 1 KB (KiloByte) = 1.000 Bytes
- 1 MB (MegaByte) = 1.000.000 Bytes
- 1 GB (GigaByte) = 1.000.000.000 Bytes

Mvήμη με N Bytes

0	0	1	0	1	0	1	0
0	1	0	0	0	0	1	0
1	1	1	0	0	0	1	1

...

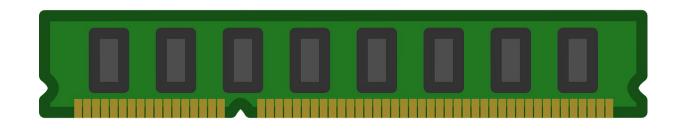
Byte 0

Byte 1

Byte 2

 Byte N-1
 0
 0
 0
 0
 0
 0
 0
 0

 Byte N
 0
 1
 1
 0
 1
 0
 0
 0



Κατηγορίες Μνήμης

Υπάρχουν 3 κατηγορίες μνήμης:

- 1. Η στοίβα (stack)
- 2. Ο σωρός (heap)
- 3. Η παγκόσμια / στατική μνήμη (global / static memory)

Η στοίβα (stack) είναι μια συνεχόμενη περιοχή της μνήμης που προστίθενται και αφαιρούνται στοιχεία με σειρά Last-In-First-Out (LIFO), δηλαδή το τελευταίο στοιχείο που προστέθηκε είναι το πρώτο που θα αφαιρεθεί.



Η στοίβα (stack) είναι μια συνεχόμενη περιοχή της μνήμης που προστίθενται και αφαιρούνται στοιχεία με σειρά Last-In-First-Out (LIFO), δηλαδή το τελευταίο στοιχείο που προστέθηκε είναι το πρώτο που θα αφαιρεθεί.

char a = 61; char b = 62; char c = 63;

Byte 1

Byte 2

Byte 3

Byte 31996

Byte 31997

Byte 31998

Byte 31999

Byte 32000

62

63

Η στοίβα (stack) είναι μια συνεχόμενη περιοχή της μνήμης που προστίθενται και αφαιρούνται στοιχεία με σειρά Last-In-First-Out (LIFO), δηλαδή το τελευταίο στοιχείο που προστέθηκε είναι το πρώτο που θα αφαιρεθεί.

```
char a = 61; char b = 62; char c = 63;
```

• • •

char d = 64;

Byte 1

Byte 2

Byte 3

Byte 31996

Byte 31997

Byte 31998

Byte 31999

Byte 32000

62

63

Η στοίβα (stack) είναι μια συνεχόμενη περιοχή της μνήμης που προστίθενται και αφαιρούνται στοιχεία με σειρά Last-In-First-Out (LIFO), δηλαδή το τελευταίο στοιχείο που προστέθηκε είναι το πρώτο που θα αφαιρεθεί.

```
char a = 61; char b = 62; char c = 63;
```

• • •

char d = 64;

Byte 1

Byte 2

Byte 3

Byte 31996

Byte 31997

Byte 31998

Byte 31999

Byte 32000

64

63

62

Η στοίβα (stack) είναι μια συνεχόμενη περιοχή της μνήμης που προστίθενται και αφαιρούνται στοιχεία με σειρά Last-In-First-Out (LIFO), δηλαδή το τελευταίο στοιχείο που προστέθηκε είναι το πρώτο που θα αφαιρεθεί.

char a = 61; char b = 62; char c = 63;

Byte 1

Byte 2

Byte 3

Byte 31996

Byte 31997

Byte 31998

Byte 31999

Byte 32000

62

63

Η στοίβα (stack) είναι μια συνεχόμενη περιοχή της μνήμης που προστίθενται και αφαιρούνται στοιχεία με σειρά Last-In-First-Out (LIFO), δηλαδή το τελευταίο στοιχείο που προστέθηκε είναι το πρώτο που θα αφαιρεθεί.

char a = 61; char b = 62;

Byte 1

Byte 2

Byte 3

Byte 31996

Byte 31997

Byte 31998

Byte 31999

Byte 32000

62

2: Ορίσματα που περνάμε στην συνάρτηση

```
int equalIgnoreCase(char char1, char char2) {

char lowerChar1 = tolower(char1);

char lowerChar2 = tolower(char2);

char lowerChar2 = tolower(char2);

return lowerChar1 == lowerChar2;

3: Προσωρινά δεδομένα
(συνήθως μερικά bytes) που
```

αποθηκεύει ο μεταγλωττιστής

2: Ορίσματα που περνάμε στην συνάρτηση

```
int equalIgnoreCase(char char1, char char2) {
```

```
1: Τοπικές Μεταβλητές ορισμένες μέσα στην συνάρτηση
```

```
char lowerChar1 = tolower(char1);
```

char lowerChar2 = tolower(char2);

return lowerChar1 == lowerChar2;

3: Προσωρινά δεδομένα (συνήθως μερικά bytes) που αποθηκεύει ο μεταγλωττιστής

CompTmp2

lowerChar2

lowerChar1

CompTmp1

char2

char1

• • •

```
int equalIgnoreCase(char char1, char char2) {
  char lowerChar1 = tolower(char1);
  char lowerChar2 = tolower(char2);
  return lowerChar1 == lowerChar2;
```

Χώρος που δεσμεύεται στην στοίβα σε **κάθε** κλήση της συνάρτησης equalignoreCase. Αυτό το κομμάτι μνήμης λέγεται και διάγραμμα ενεργοποίησης (activation record ή stack frame) της συνάρτησης

CompTmp2

lowerChar2

lowerChar1

CompTmp1

char2

char1

• • •

```
char tolower(char c) {
    if (c >= 'A' && c <= 'Z')
        return c + ('a' - 'A');
    return c;
int equalIgnoreCase(char char1, char char2) {
 char lowerChar1 = tolower(char1);
                                                equalIgnoreCase
 char lowerChar2 = tolower(char2);
 return lowerChar1 == lowerChar2;
```

Τι θα συμβεί όταν κληθεί η συνάρτηση tolower την πρώτη φορά;

CompTmp2 lowerChar2 lowerChar1 CompTmp1 char2 char1

```
char tolower(char c) {
                                                       tolower
    if (c >= 'A' && c <= 'Z')
        return c + ('a' - 'A');
    return c;
int equalIgnoreCase(char char1, char char2) {
 char lowerChar1 = tolower(char1);
                                                equalIgnoreCase
 char lowerChar2 = tolower(char2);
 return lowerChar1 == lowerChar2;
```

CompTmp5

CompTmp4

CompTmp3

C

CompTmp2

lowerChar2

lowerChar1

CompTmp1

char2

char1

. . .

Τι θα συμβεί όταν εκτελεστεί η εντολή return;

```
char tolower(char c) {
                                                       tolower
    if (c >= 'A' && c <= 'Z')
        return c + ('a' - 'A');
   return c;
int equalIgnoreCase(char char1, char char2) {
 char lowerChar1 = tolower(char1);
                                                equalIgnoreCase
 char lowerChar2 = tolower(char2);
 return lowerChar1 == lowerChar2;
```

CompTmp5

CompTmp4

CompTmp3

C

CompTmp2

lowerChar2

lowerChar1

CompTmp1

char2

char1

. . .

```
char tolower(char c) {
    if (c >= 'A' && c <= 'Z')
        return c + ('a' - 'A');
    return c;
int equalIgnoreCase(char char1, char char2) {
 char lowerChar1 = tolower(char1);
                                                equalIgnoreCase
 char lowerChar2 = tolower(char2);
 return lowerChar1 == lowerChar2;
```

CompTmp2 lowerChar2 lowerChar1 CompTmp1 char2 char1

Τι θα συμβεί όταν εκτελεστεί η δεύτερη κλήση tolower;

Τι θα συμβεί όταν εκτελεστεί η εντολή return;

```
char tolower(char c) {
                                                       tolower
    if (c >= 'A' && c <= 'Z')
        return c + ('a' - 'A');
   return c;
int equalIgnoreCase(char char1, char char2) {
 char lowerChar1 = tolower(char1);
                                                equalIgnoreCase
char lowerChar2 = tolower(char2);
 return lowerChar1 == lowerChar2;
```

CompTmp5

CompTmp4

CompTmp3

C

CompTmp2

lowerChar2

lowerChar1

CompTmp1

char2

char1

. . .

```
char tolower(char c) {
    if (c >= 'A' && c <= 'Z')
        return c + ('a' - 'A');
    return c;
int equalIgnoreCase(char char1, char char2) {
 char lowerChar1 = tolower(char1);
                                                equalIgnoreCase
 char lowerChar2 = tolower(char2);
 return lowerChar1 == lowerChar2;
```

Τι θα συμβεί όταν εκτελεστεί η return της equalignoreCase;

CompTmp2

lowerChar2

lowerChar1

CompTmp1

char2

char1

. . .

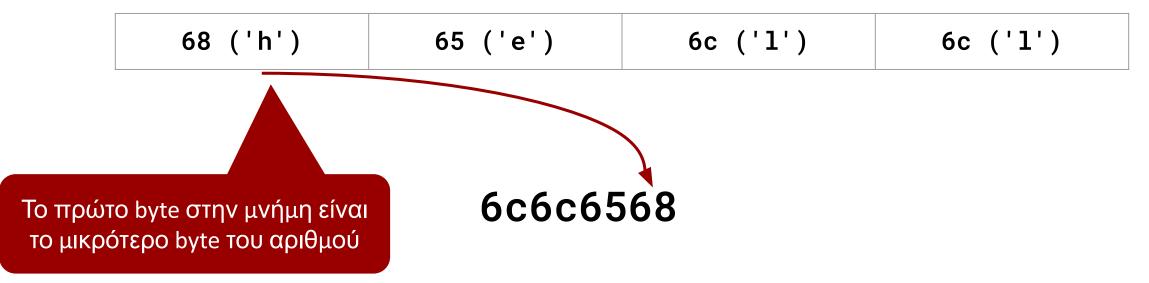
```
char tolower(char c) {
    if (c >= 'A' && c <= 'Z')
        return c + ('a' - 'A');
    return c;
int equalIgnoreCase(char char1, char char2) {
 char lowerChar1 = tolower(char1);
 char lowerChar2 = tolower(char2);
 return lowerChar1 == lowerChar2;
```

Τι θα συμβεί όταν εκτελεστεί η return της equalIgnoreCase;

Πόσα bytes έχει ένας int;

Endianness

Endianness λέμε τον τρόπο με τον οποίο οι ακέραιοι αποθηκεύονται στην μνήμη. Οι ακέραιοι αποτελούνται από πολλά bytes και επομένως πρέπει να αποφασίσουμε αν τους αποθηκεύουμε από το μικρότερο στο μεγαλύτερο (little endian) ή από το μεγαλύτερο στο μικρότερο (big endian).



Παράδειγμα Endianness

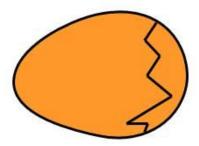
Τι θα τυπώσει το παρακάτω:

```
#include <stdio.h>
int main() {
  int x = 0x42434445;
  char * bytes = (char*)&x;
  int i;
  for(i = 0; i < sizeof(int) / sizeof(char); i++)</pre>
  printf("%02x\n", bytes[i]);
  return 0;
```

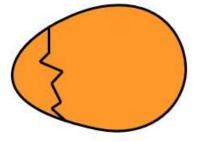
Παράδειγμα Endianness

```
$ ./int
                                                          45
Τι θα τυπώσει το παρακάτω:
                                                          44
                                                          43
     #include <stdio.h>
                                                          42
     int main() {
       int x = 0x42434445;
       char * bytes = (char*)&x;
       int i;
       for(i = 0; i < sizeof(int) / sizeof(char); i++)</pre>
        printf("%02x\n", bytes[i]);
       return 0;
```

ENDIANNESS



BIG ENDIAN - The way people always broke their eggs in the Lilliput land



LITTLE ENDIAN - The way the king then ordered the people to break their eggs

- <u>Little endian</u> (x86/x86_64/arm/...): LSB stored in smallest memory address
- Big endian (PPC/SPARC/...):
 MSB stored in smallest memory address
- <u>Bi-endian</u> (Some arm and mips):
 Can switch dynamically

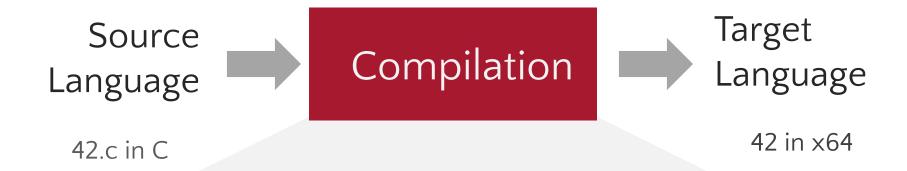
Compilation
Workflow &
Execution
Semantics

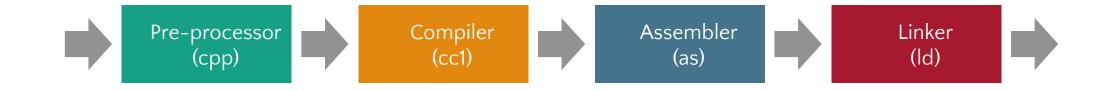
To answer

Is this program safe?

We must first know

What will executing it do?





Preprocessor

Compiler

Assembler

> Linker

\$ gcc -E

```
#include <stdio.h>

void answer_function(char *name, int x)
{
  printf("% s answered % d\n", name, x);
}

int main(int argc, char *argv[])
{
  answer_function("hello", 42);
}
```

#include expansion #define substitution

Preprocessor

Compiler

Assembler

Linker

```
$ gcc -S
```

```
answer_function:
.LFB0:
.cfi_startproc
pushq %rbp
.cfi_def_cfa_offset 16
.cfi_offset 6, -16
movq %rsp, %rbp
.cfi_def_cfa_register 6
subq $16, %rsp
movq %rdi, -8(%rbp)
movl %esi, -12(%rbp)
movl -12(\%rbp), %edx
```

Generate assembly code



Compiler

Assembler

> Linker

```
$ as hello.s -o hello.o
$ file hello.o
hello.o: ELF 64-bit LSB relocatable,
x86-64, version 1 (SYSV), not stripped
```

Create object code

(the ones and zeros a computer executes)

Preprocessor

Compiler

Assembler

Linker

\$ ld <opts> (this one is complicated; see gcc -v)

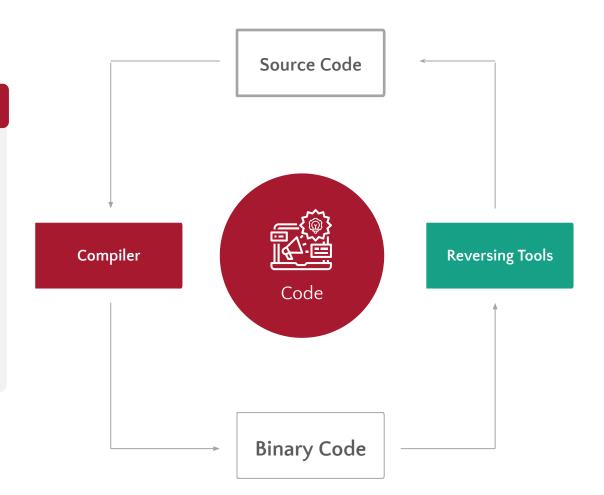
```
gcc -v hello.c
/usr/lib/gcc/x86 64-linux-gnu/10/collect2 -plugin
/usr/lib/gcc/x86 64-linux-gnu/10/liblto plugin.so
-plugin-opt=/usr/lib/gcc/x86_64-linux-gnu/10/lto-wrapper
-plugin-opt=-fresolution=/tmp/ccLvq7Mj.res -plugin-opt=-pass-through=-lgcc
-plugin-opt=-pass-through=-lgcc s -plugin-opt=-pass-through=-lc
-plugin-opt=-pass-through=-lgcc -plugin-opt=-pass-through=-lgcc s --build-id --eh-frame-hdr
-m elf x86 64 --hash-style=gnu --as-needed -dynamic-linker /lib64/ld-linux-x86-64.so.2 -pie
/usr/lib/gcc/x86_64-linux-gnu/10/../../x86_64-linux-gnu/Scrt1.o
/usr/lib/gcc/x86 64-linux-gnu/10/\dots/\dots/x86 64-linux-gnu/crti.o
/usr/lib/gcc/x86 64-linux-gnu/10/crtbeginS.o -L/usr/lib/gcc/x86 64-linux-gnu/10
-L/usr/lib/gcc/x86 64-linux-gnu/10/../../x86 64-linux-gnu
-L/usr/lib/gcc/x86 64-linux-gnu/10/../../../lib -L/lib/x86 64-linux-gnu -L/lib/../lib
-L/usr/lib/x86 64-linux-gnu -L/usr/lib/../lib -L/usr/lib/gcc/x86 64-linux-gnu/10/../../..
/tmp/ccRnv5wm.o -lgcc --push-state --as-needed -lgcc s --pop-state -lc -lgcc --push-state
/usr/lib/gcc/x86 64-linux-gnu/10/../../x86 64-linux-gnu/crtn.o
```

Linker links with other compilation units (.o) and libraries (.a) to produce an executable.

Compilation is not necessarily invertible

Compilation

Turn source code into executable code using programmer-defined abstractions as a guide translation.



Reversing

Best case effort to recover source code from binary code.

Interview Question: How would you see the content of a binary program?

Binary

Code Segment (.text)

Data Segment (.data)

...

The final executable binary

- Text segment
- Data for constants like "hello world" and globals
- And more...



```
# read the various sections. Also see `nm`
$ readelf -S <file>

# -d: disassemble
# -S: match symbol locations to source
$ objdump <-d> <-S> file
```

BASIC EXECUTION MODEL

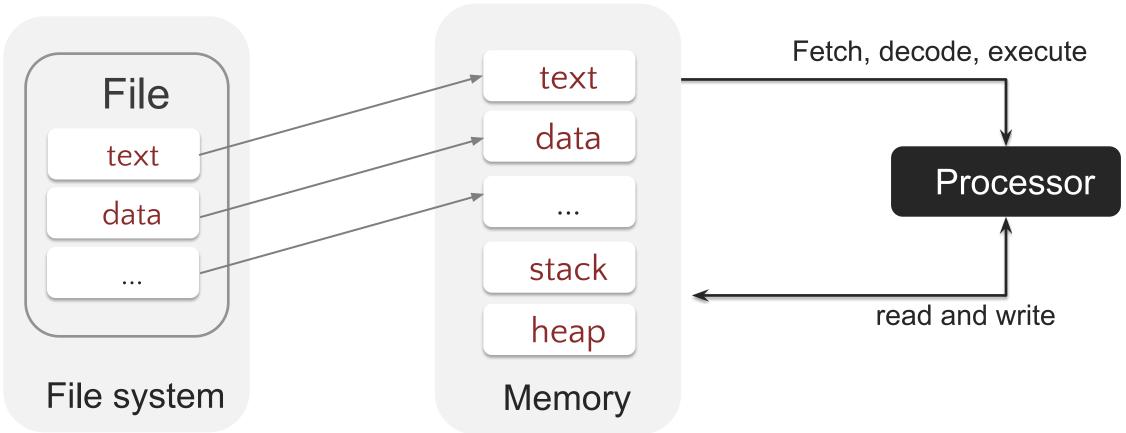


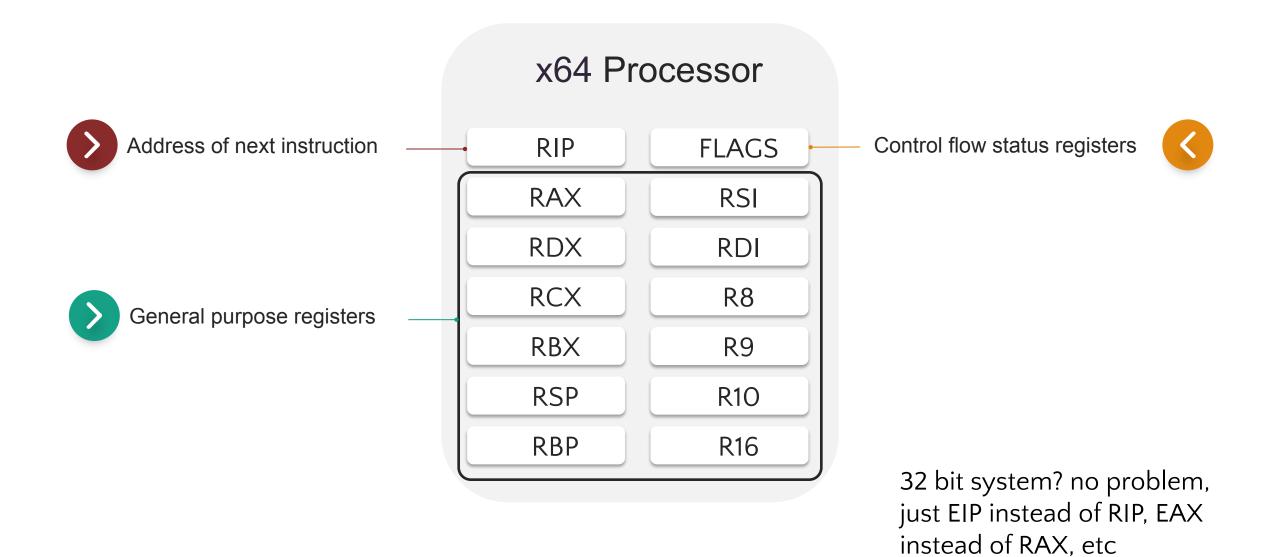
2

Segments loaded and process memory created

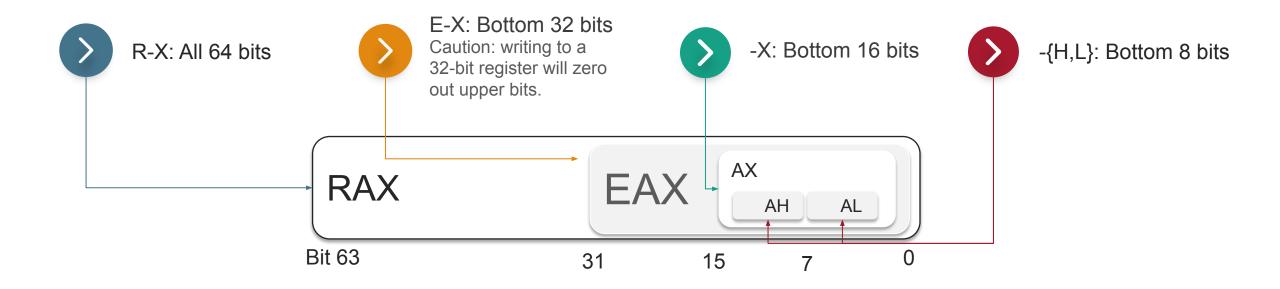
3

Processor executes process in memory





Parts of a register can be addressed directly



Basic Ops and AT&T Vs. Intel Syntax

Meaning	AT&T	Intel
rbx = rax	movq %rax, %rbx	mov rbx, rax
rax = rax + rbx	addq %rbx, %rax	add rax, rbx
$rex = rex \ll 2$	shl \$2, %rcx	shl rex, 2



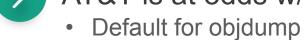


r

Intel mirrors assignment order

- objdump –M intel
- Default for windows





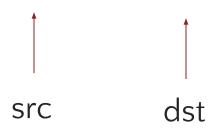
Default in UNIX

Memory Operations

Memory Loads

What do these instructions do?

- 1. movb (%rax), dl
- 2. mov (%rax), edx



- Bracket "[]" indicates Intel
- Paren "()" indicates AT&T

Register	Value
rax	0x3
rdx	0x0
rbx	0x5

Byte	Address
0xff	6
0xee	5
0xdd	4
0xcc	3
0xbb	2
0xaa	1
0x00	0

Memory Loads

What do these instructions do?

- 1. movb (%rax), dl
- 1. mov (%rax), edx

Register	Value
rax	0x3
rdx	0x0
rbx	0x5

Byte	Address
0xff	6
0xee	5
0xdd	4
0xcc	3
0xbb	2
0xaa	1
0x00	0

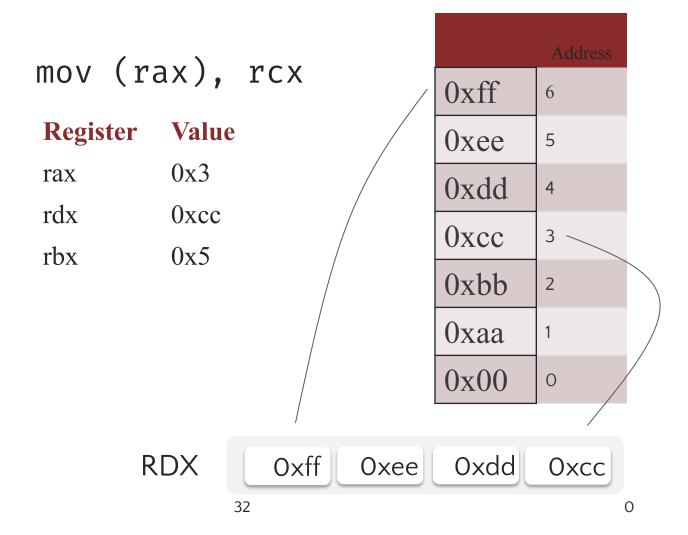


- 1. Moves 0xcc into dl
- 2. Moves 0xcc 0xff into edx. But what number is this?

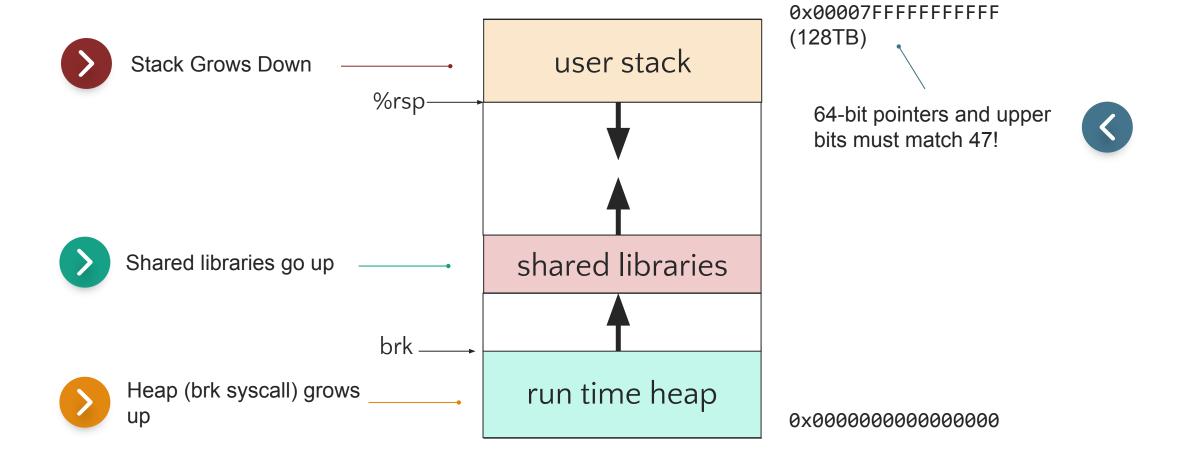
EXAMPLE

We'll assume x86/x86_64 from now on:

- Address a goes into register bits 0-7
- Address a+1 in the 8-15
- And so on



Process Memory Organization

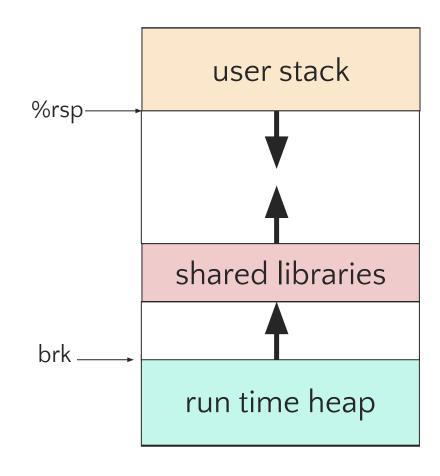


On the stack

- Local variables
- Lifetime:stack frame

On the heap

- Dynamically allocated via new/malloc/etc.
- Lifetime: until freed

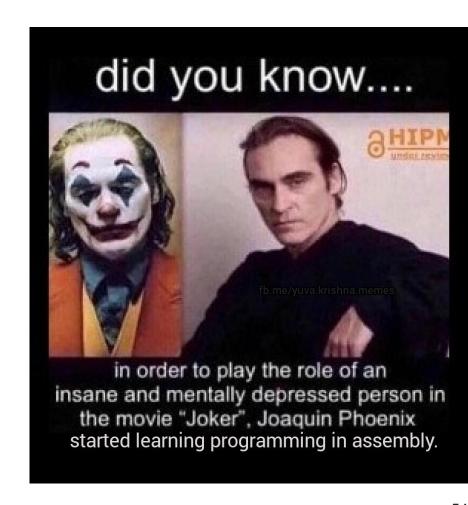


0x00007FFFFFFFFFF (128TB)

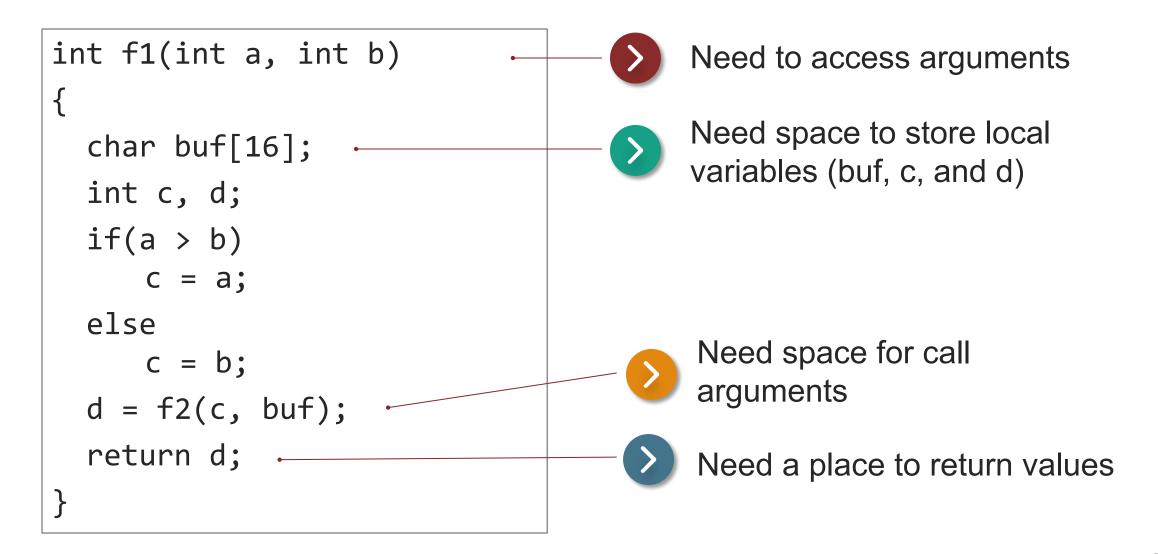
0x0000000000000000

Procedures

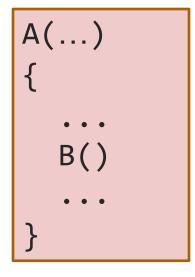
- Assembly code doesn't have procedures
- Compilers implement procedures
 - On the stack
 - Following the call/return stack discipline

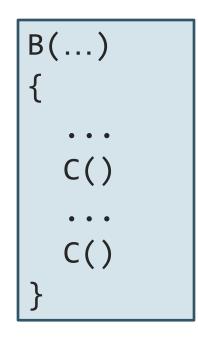


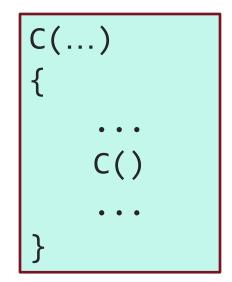
The compiler needs to implement a call stack, calling convention, and a way to pass arguments for code being compiled, third-party binary libraries, and the OS



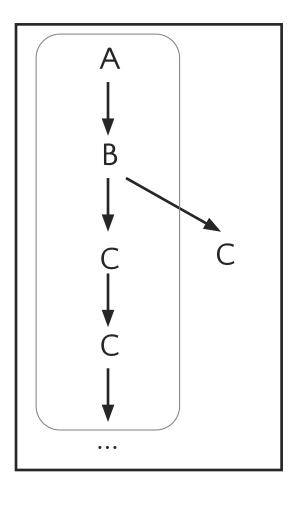
Recall the basic stack semantics of function calls



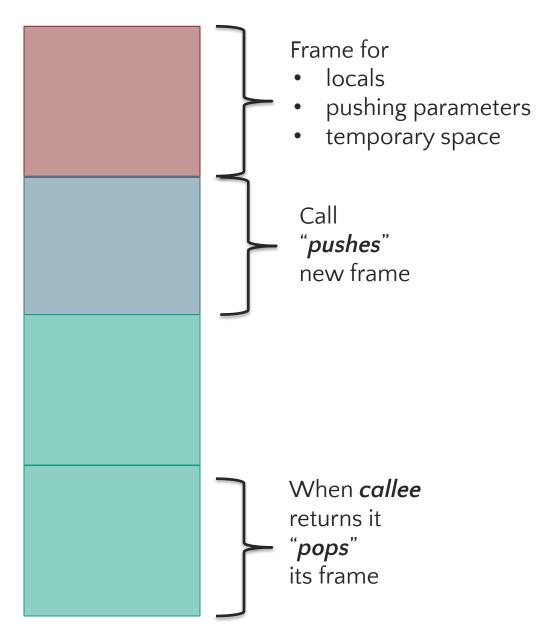




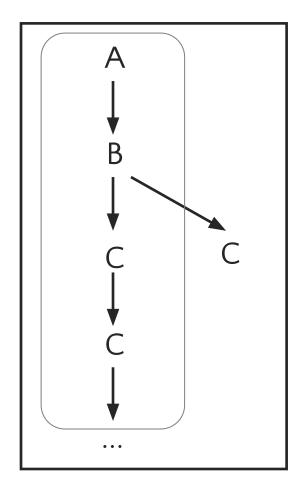
Function Call Chain



The Stack



Function Call Chain



A Toy Example

```
void print(char * message) {
   char buffer[64];
   strcpy(buffer, message);
   printf("%s\n", buffer);
int main(int argc, char ** argv) {
   print(argv[1]);
   return 0;
```

```
080491d7 <print>:
                                         . . .
A Toy Example
                                         80491f2:
                                                        e8 69 fe ff ff
                                                                         call 8049060 <strcpy@plt>
                                         80491f7:
                                                                         add $0x10, %esp
                                                        83 c4 10
void print(char * message) {
    char buffer[64];
                                         8049202:
                                                        e8 69 fe ff ff
                                                                         call 8049070 <puts@plt>
    strcpy(buffer, message);
                                         8049207:
                                                                          add $0x10, %esp
                                                        83 c4 10
    printf("%s\n", buffer);
                                         804920f:
                                                        c3
                                                                          ret
int main(int argc, char ** argv) {
                                                                         Return to the caller using
                                        08049210 <main>:
                                                                           the ret instruction
    . . .
    print(argv[1]);
                                         804925e:
                                                        e8 74 ff ff ff
                                                                         call 80491d7 <print>
    return 0;
                                         8049263:
                                                        83 c4 10
                                                                         add $0x10, %esp
          Transfer control to functions
                                                                                                 59
          using the call instruction
                                         8049274:
                                                        c3
                                                                         ret
```

Call - Ret Semantics

Call semantics: Store address of next instruction (the return address) where the current stack pointer is pointing and then jump to target address. call target is equivalent to:

push next_address
jmp target

mem[sp] = next_addr; sp -= sizeof(next_addr)

Ret semantics: Read address of next instruction from the stack (the return address) and jump to it. ret is equivalent to:

pop next_address
jmp next_address

next_addr = mem[sp]; sp += sizeof(next_addr)

```
0xffffd66c
A Toy Example
                                                                   main return addr
void print(char * message) {
                                                                   main locals vars
   char buffer[64];
                                                     0xffffd640
   strcpy(buffer, message);
   printf("%s\n", buffer);
int main(int argc, char ** argv) {
   . . .
   print(argv[1]);
   return 0;
           804925e: e8 74 ff ff ff call 80491d7 <print>
           esp = 0xffffd640; eip = 0x804925e
```

61

```
0xffffd66c
A Toy Example
                                                                  main return addr
void print(char * message) {
                                                                  main locals vars
   char buffer[64];
   strcpy(buffer, message);
                                                                 print return addr
                                                    0xffffd640
                                                                      0x8049263
   printf("%s\n", buffer);
int main(int argc, char ** argv) {
   . . .
   print(argv[1]);
   return 0;
          80491d7: 53
                              push
                                    %ebx
                                                                                     62
          esp = 0xffffd63c; eip = 0x80491d7
```

```
0xffffd66c
A Toy Example
                                                                 main return addr
void print(char * message) {
                                                                 main locals vars
   char buffer[64];
   strcpy(buffer, message);
                                                                 print return addr
                                                    0xffffd640
                                                                     0x8049263
   printf("%s\n", buffer);
                                                                    print locals
                                                                  including buffer
int main(int argc, char ** argv) {
                                                                    (at least 64
                                                                       bytes)
   . . .
                                                    0xffffd5e0
   print(argv[1]);
   return 0;
     8049202: e8 69 fe ff ff call 8049070 <puts@plt>
     esp = 0xffffd5e0; eip = 0x8049202
```

```
0xffffd66c
A Toy Example
                                                                   main return addr
void print(char * message) {
                                                                   main locals vars
   char buffer[64];
   strcpy(buffer, message);
                                                                  print return addr
                                                     0xffffd640
                                                                       0x8049263
   printf("%s\n", buffer);
int main(int argc, char ** argv) {
   . . .
                                                     0xffffd5e0
   print(argv[1]);
   return 0;
     804920f:
               c3
                                           ret
     esp = 0xffffd63c; eip = 0x804920f
```

```
0xffffd66c
A Toy Example
                                                                   main return addr
void print(char * message) {
                                                                   main locals vars
   char buffer[64];
   strcpy(buffer, message);
                                                     0xffffd640
   printf("%s\n", buffer);
int main(int argc, char ** argv) {
   . . .
                                                      0xffffd5e0
   print(argv[1]);
   return 0;
     8049263: 83 c4 10
                                           add $0x10, %esp
     esp = 0xffffd640; eip = 0x8049263
```

Buffer Overflows

What Are Buffer Overflows?

A *buffer overflow* occurs when data is written <u>outside</u> of the space allocated for the buffer

C does not check that writes are in-bounds

1. Stack-based

- what we'll cover today
- 2. Heap-based
 - more advanced
 - very dependent on system and library version

```
0xffffd66c
                        Input: ./bof0 "ABCD"
A Toy Example
                                                                  main return addr
void print(char * message) {
                                                                  main locals vars
   char buffer[64];
   strcpy(buffer, message);
                                                                  print return addr
                                                    0xffffd640
                                                                      0x8049263
   printf("%s\n", buffer);
                                                                          ABCD
int main(int argc, char ** argv) {
   . . .
                                                     0xffffd5e0
   print(argv[1]);
   return 0;
     8049202: e8 69 fe ff ff call 8049070 <puts@plt>
```

esp = 0xffffd5e0; eip = 0x8049202

```
0xffffd66c
                    Input: ./bof0 "AAAAAA...AAA"
A Toy Example
                                                                   main return addr
void print(char * message) {
                                                                   main locals vars
   char buffer[64];
   strcpy(buffer, message);
                                                     0xffffd640
   printf("%s\n", buffer);
int main(int argc, char ** argv) {
   . . .
                                                     0xffffd5e0
   print(argv[1]);
   return 0;
     8049202: e8 69 fe ff ff call 8049070 <puts@plt>
     esp = 0xffffd5e0; eip = 0x8049202
```

```
0xffffd66c
                  Input: ./bof0 "AAAAAA...AAAAA...AAA"
A Toy Example
void print(char * message) {
   char buffer[64];
   strcpy(buffer, message);
                                                    0xffffd640
   printf("%s\n", buffer);
int main(int argc, char ** argv) {
                                                     0xffffd5e0
   print(argv[1]);
   return 0;
     8049202: e8 69 fe ff ff call 8049070 <puts@plt>
```

esp = 0xffffd5e0; eip = 0x8049202

```
0xffffd66c
                  Input: ./bof0 "AAAAAA...AAAA"
A Toy Example
void print(char * message) {
   char buffer[64];
   strcpy(buffer, message);
                                                    0xffffd640
   printf("%s\n", buffer);
int main(int argc, char ** argv) {
                                                    0xffffd5e0
   print(argv[1]);
   return 0;
     804920f:
               c3
                                          ret
```

esp = 0xffffd63c; eip = 0x804920f

```
0xffffd66c
                  Input: ./bof0 "AAAAAA...AAAAA...AAA"
A Toy Example
void print(char * message) {
   char buffer[64];
   strcpy(buffer, message);
                                                     0xffffd640
   printf("%s\n", buffer);
int main(int argc, char ** argv) {
                                                      0xffffd5e0
   print(arg)
             Program received signal SIGSEGV, Segmentation fault.
   return 0;
              0x41414141 in ?? ()
              (gdb) i r eip
                            0x41414141
                                             0x41414141
              eip
```

```
0xffffd66c
                  Input: ./bof0 "AAAAAA...AAAAA...AAA"
A Toy Example
void print(char * message) {
   char buffer[64];
   strcpy(buffer, message);
                                                    0xffffd640
   printf("%s\n", buffer);
int main(int argc, char ** argv) {
                                                     0xffffd5e0
   print(argv[1]);
   return 0 $ dmesg -T | grep segfault
             [Thu Mar 21 10:58:52 2024] bof0_real[21893]: segfault at
             41414141 ip 0000000041414141 sp 0000000ffffd640 error 14
             in libc.so.6[f7d85000+20000]
```

Can we return control to another function?

Let's do it Live!

Shellcode Injection

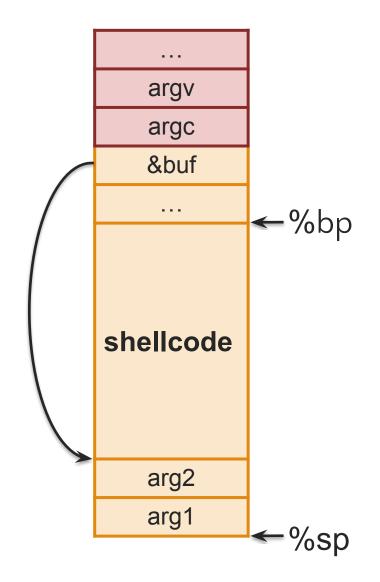
Shellcode

Traditionally exploits injected assembly instructions for exec("/bin/sh") into buffer.

Data Execution Prevention and other defenses have made this exploitation technique ineffective on consumer commercial OSes for over a decade.

Sadly, this is still applicable in areas like IoT, energy, and so on.

- Considered a basic skill for exploitation (even if not on your latest OS)
- See "Smashing the stack for fun and profit" for one string
- or search online OR write it yourself!

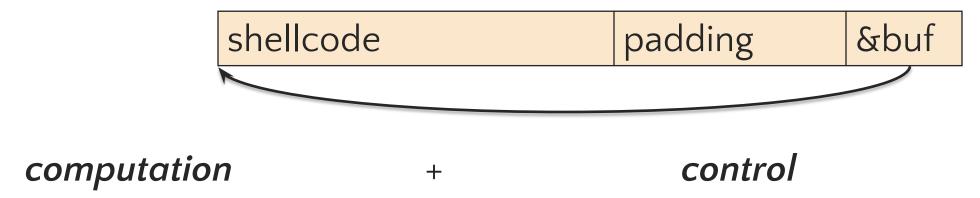


Recap

To generate *exploit* for a basic buffer overflow:

Determine size of stack frame up to head of buffer

1. Overflow buffer with the right size



Ευχαριστώ και καλή μέρα εύχομαι!

Keep hacking!